

Fig. 1 No free lunch: model behaviour is dataset dependent





Fig. 2 Cost and accuracy span two orders of magnitude

(a) Accuracy vs cost (marker size = epochs)



(b) Best achievable accuracy under a budget





Fig. 3 Routing recovers most of the oracle gain

(a) Router vs baselines vs oracle



(b) Per-dataset gap (6-pool, single-best=LOCF)





Fig. 4 Statistical comparison of the model pool

(a) Mean ranks $\pm$ CD/2 (Friedman $p = 1.2e - 66$)



(b) Wilcoxon vs single-best (GPT4TS)





Fig. 5 Robustness to the missing rate

(a) Per-model error vs missing rate



(b) Routing vs baselines (6-pool)





Fig. 6 Which C-TCAR features the router relies on

(a) Top-12 individual features



(b) Importance by feature group





Fig. 7 The cross-dataset router overfits its feature set

(a) Remove one feature group (LOGO)



(b) Robust in-pool, overfits cross-dataset





Fig. 8 Router selection quality

(a) Generalization gap: LOO vs LOGO



(b) Routed vs oracle (6-pool, LOGO, regret 0.029)





Fig. 9 Where cross-dataset routing errs and what it costs

(a) Accuracy per held-out dataset (6-pool)



(b) Regret: LOO vs LOGO (6-pool)





Fig. 10 Training dynamics of the deep imputers

(a) Convergence (marker = 90% of total drop)



(b) Convergence speed

